Live Demo Materials: Varity (Final, 1 Nov 2026)

Team Hokkien Mee Is Purple · 10 min pitch + live demo, then 5 min Q&A (−1 point per 30 s over). The judges watch one synthetic user go through all six stages in one live session.

1. What judges need

Item	Where
Live app	https://verity--eezhan0805.replit.app
Code	https://github.com/TayEeZhan/varity (public)
Synthetic CVs to upload	`demo-cvs/` in the repo: `aisyah-backend.docx`, `arjun-data.odt`, `daniel-devops.txt` / `.md` / `.html`, `sitihajar-ml.rtf`, `weijie-frontend.pdf` (all synthetic). In the app: Aina (her CV is read, with quotes), Ben and Chloe (skills only)
Login	Judges sign in with their own Replit or Google account ("Get started"). No test password is shared. On the demo laptop: the team's demo account, already signed in to Google in a normal window (not incognito)
Backup	Local copy with login off (section 5), and the backup video [TEAM TO FILL: YouTube link, if recorded]
Proof files	`data/impact_check.json` (impact run), `data/backtest.json`, `data/role_backtest.json`, `data/llm_usage.csv`, the GitHub Actions run log

2. Before the slot (15 minutes before)

- [ ] Open the live app once to wake it (Autoscale sleeps; the first load can take about 30 s).
- [ ] GitHub → Actions → "Daily data refresh" → Run workflow. Wait for the green tick; Job forecast then says "updated ".
- [ ] Demo account: account menu (top right) → Delete account…, so the demo starts as a new user.
- [ ] Open a second tab on `#/radar` for the replay check.
- [ ] Local backup running on the laptop (section 5); HDMI adapter; charger.
- [ ] Open every link once in an incognito window.

3. The six-stage journey (about 6 minutes)

#	Stage (guide wording)	Do this	Show and say
1	**Tell us about yourself**	Landing → **Get started** → sign in → pop-up "Start with your CV" → tick the synthetic box → drop `aisyah-backend.docx` (or pick **Aina (synthetic)**)	Each skill appears with **the sentence from the CV that proves it**. Untick one skill to show the user controls the profile, then press **Save N skills**. "Skills come from the CV; nothing is pre-filled."
2	**Future Skills Radar**	Landing → **Open the skills radar** (`#/radar`); then the **Signals** and **Job forecast** tabs	Rising and declining skills, each with **source and date** ("Source: HN Who is hiring · as of 2026-10 · fetched …"). Signals: AI news ranked from dated sources, with scores. Job forecast: roles with range bars, the Past year change and Growing / Steady / Shrinking.
3	**Get matched**	**CV & routes** → open **Why NN%?** on the top route	Ranked roles, each with a **match %** and **skills gap**. The breakdown shows 0.7 × skill overlap + 0.3 × outlook, and **What would raise it** lists each gap skill with the new score and a free course.
4	**Follow a pathway**	**Pathway** tab (the top route is the default; **Switch to this route** picks another), then **Today**	Ordered stages with **minutes per step** and free resources (click one: the real course opens), plus an **HRD Corp-claimable** option with days and fee. This week is open; later weeks are locked until it is done.
5	**SDG contribution**	Pathway → the SDG section; **I care about** (SDG goal picker)	Every role and skill has an SDG target **with a one-line reason**, and the plan adds them up ("How this plan contributes: SDG 8.2 (4), SDG 7.3 (1)"). Pick a goal: the route order changes, so the tag changes a choice.
6	**Compare futures**	Toggle **Expected (estimate)** ↔ **If AI speeds up (estimate)** on CV & routes or Pathway	Scores and the plan shift; both are labelled estimates. "The scenario applies one explicit parameter, k = 2, to each skill's AI exposure; it is not a hand-written alternative."

4. Live checks the judges may ask for

Judge asks	Do this	What to say
"Change the trend data now."	Skills radar → **Replay** → pick 2023-12 → open CV & routes and Pathway	The radar, the role scores and the plan change, because they read that month's data. **Back to latest** restores it. **Fetch latest HN now** pulls the newest posts (one fetch per 10 minutes).
"Is this course / number real?"	Click a resource link in the Pathway; click an HN thread link on Job forecast	Every resource link was checked (HTTP 200); every number names its file and date.

Judge asks	Do this	What to say
"Why this score? What would raise it?"	**Why NN%?** on the route	match = Σ weights held ÷ Σ role weights; rank = 0.7 × match + 0.3 × outlook; each gap skill shows the new score.
"How do you keep data current? Cost?"	Job forecast header "updated …"; GitHub Actions log	Daily at 09:00 MYT. Labelling all 48,491 posts cost US$2.63 once; a CV read costs at most US$0.0041.
"What if AI moves faster or slower?"	Scenario toggle; Method screen	k = 2 is an assumption; the data give −1.0 to 2.9; the Fast scenario shows the effect, labelled as an estimate.
"How do you stop it inventing?"	Upload a recipe file → "This does not look like a CV … try again"; the review line "Dropped because the quote is not in your CV"	Models only point at quotes; code checks each quote and sets every number, rank and course.
"Is it fair?"	Method → **Fairness check**	The test swaps name, age, gender, nationality and school; the ranking does not move.

5. Backup if Wi-Fi or Replit fails

On the demo laptop (login is off on a local copy, so no sign-in is needed), from the `verity` folder:

```
cd backend\signals
..\..\.venv\Scripts\python -m trendsignal serve --host 127.0.0.1 --port 8000
```

Open http://127.0.0.1:8000 and run the same six stages. If the laptop also fails, play the backup video.

6. Say this honestly

- Predictions are estimates: we show an 80% band, and "Steady" when the change could be noise.
- Hacker News is a US and startup sample; DOSM gives a Malaysian sector check, not a per-role number.
- We tested with synthetic profiles only; the three Phase 1 user KPIs need a real-user pilot.
- Varity supports a decision; it does not make hiring decisions.